

Module 05 Lab - Data Preparation Objective: To learn and apply the most common data preparation techniques. Raw data is rarely ready for a machine learning model. This process, also called preprocessing, is one of the most critical steps in the entire ML workflow. **In this lab, you will write more of the code.** Read the explanations and then complete the tasks in the code cells.

Part 1: Setup and Initial Look We will continue using the Titanic dataset because it has the exact problems we need to solve: missing values and non-numeric data.

```
In [7]: import pandas as pd
import numpy as np # Load the dataset
df = pd.read_csv('https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv') #
print("--- Missing Values Before Cleaning ---")
print(df.isnull().sum())
```

```

--- Missing Values Before Cleaning ---
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin           687
Embarked         2
dtype: int64

```

Part 2: Handling Missing Values (Imputation)
Concept: Most machine learning models cannot handle missing values (NaN). We must deal with them. Dropping the rows is an option, but you lose data. A better way is **imputation**, which means filling in the missing values with a calculated guess. Common imputation strategies:
 * **Mean:** Fill with the average value. Good for normally distributed data.
 * **Median:** Fill with the middle value. Better for skewed data or data with outliers (like Fare).
 * **Mode:** Fill with the most frequent value. Used for categorical data.

Task 1: Impute the 'Age' Column
 The 'Age' column is missing many values. Since age can be skewed (e.g., by a few very old passengers), using the

median is a robust choice. **Your Task:** Calculate the median of the 'Age' column and use the `.fillna()` method to replace the missing values.

```
In [8]: # --- ENTER YOUR CODE HERE ---#  
#1. Calculate the median of the 'Age' column  
median_age = df['Age'].median()  
  
# 2. Fill the missing values in 'Age' with the median#  
df['Age'].fillna(median_age, inplace=True)  
  
# 3. Verify that there are no more missing values in 'Age'  
print("Missing values in 'Age' after imputation:")  
print(df['Age'].isnull().sum())
```

Missing values in 'Age' after imputation:

0

C:\Users\safab\AppData\Local\Temp\ipykernel_27988\746889604.py:6: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or 'df[col] = df[col].method(value)' instead, to perform the operation inplace on the original object.

```
df['Age'].fillna(median_age, inplace=True)
```

Part 3: Encoding Categorical Features
Concept: Machine learning models are mathematical, so they need numbers, not text. We need to convert categorical columns (like 'Sex' and 'Embarked') into a numerical format. The most common method is **One-Hot Encoding**. One-Hot Encoding takes a

column with N categories and turns it into N new columns, each with a 1 or 0. For example, the 'Sex' column (male, female) becomes two new columns: Sex_male and Sex_female. Pandas has a convenient function called `pd.get_dummies()` that does this for us.

Task 2: One-Hot Encode Categorical Columns
Your Task: Use `pd.get_dummies()` to encode the 'Sex' and 'Embarked' columns. Make sure to drop the original columns after encoding.

```
In [9]: # --- ENTER YOUR CODE HERE ---
# 1. Use get_dummies to create new columns for 'Sex' and 'Embarked'
# Set `drop_first=True` to avoid multicollinearity (a statistical issue), which drops one of the new
encoded_df = pd.get_dummies(
    df,
    columns=['Sex', 'Embarked'],
    drop_first=True)
# 2. Display the first few rows of the new DataFrame to see the new columns
print(encoded_df.head())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Age	SibSp	Parch	\
0	Braund, Mr. Owen Harris	22.0	1	0	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	38.0	1	0	
2	Heikkinen, Miss. Laina	26.0	0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	35.0	1	0	
4	Allen, Mr. William Henry	35.0	0	0	

	Ticket	Fare	Cabin	Sex_male	Embarked_Q	Embarked_S
0	A/5 21171	7.2500	NaN	True	False	True
1	PC 17599	71.2833	C85	False	False	False
2	STON/O2. 3101282	7.9250	NaN	False	False	True
3	113803	53.1000	C123	False	False	True
4	373450	8.0500	NaN	True	False	True

Part 4: Feature Scaling**Concept:** Many models are sensitive to the scale of the features. For example, **Age** (from 0-80) and **Fare** (from 0-512) are on very different scales. This can cause the model to incorrectly believe that **Fare** is a more important feature simply because its values are larger.**Feature Scaling** solves this by putting all features on a similar scale. A common method is **Standardization** (**StandardScaler** in scikit-learn), which rescales the data to have a mean of 0 and a standard deviation of 1.**Important:** You only scale your numerical

features, not your target variable or your newly encoded categorical columns.

Task 3: Scale the 'Age' and 'Fare' Columns Your Task: Use `StandardScaler` from `sklearn.preprocessing` to scale the 'Age' and 'Fare' columns.

```
In [10]: from sklearn.preprocessing import StandardScaler
# --- ENTER YOUR CODE HERE ---
# 1. Create an instance of the StandardScaler
scaler = StandardScaler()

# 2. Select the columns to scale
columns_to_scale = ['Age', 'Fare']

# 3. Fit the scaler to the data and transform it
# Note: We are using the `encoded_df` from the previous step if you created it.
encoded_df[columns_to_scale] = scaler.fit_transform(
    encoded_df[columns_to_scale])

# 4. Display the first few rows to see the scaled data
print(encoded_df.head())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Age	SibSp	Parch	\
0	Braund, Mr. Owen Harris	-0.565736	1	0	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	0.663861	1	0	
2	Heikkinen, Miss. Laina	-0.258337	0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	0.433312	1	0	
4	Allen, Mr. William Henry	0.433312	0	0	

	Ticket	Fare	Cabin	Sex_male	Embarked_Q	Embarked_S
0	A/5 21171	-0.502445	NaN	True	False	True
1	PC 17599	0.786845	C85	False	False	False
2	STON/O2. 3101282	-0.488854	NaN	False	False	True
3	113803	0.420730	C123	False	False	True
4	373450	-0.486337	NaN	True	False	True

 **Knowledge Check Instructions:** Answer the following questions in this markdown cell.

1. Why is it often better to impute missing values with the median instead of the mean?
2. Explain in your own words what One-Hot Encoding does and why it is necessary.
3. Would you need to apply Feature Scaling to a Decision Tree model? Why or why not? (Hint: Think about how a Decision Tree makes its splits).

1. The median is often better because it is not greatly affected by extremely high and low values. It also gives a better middle value for skewed data, while the mean can be changed by outliers. 2. One-Hot Encoding changes categorical text values to numerical columns containing True/False or 1/0. This tool is

necessary because ML models needs only numerical data instead of text. 3. No, Feature Scalling is not needed for a Decision Tree. A Decision Tree makes splits by comparing values to thresholds, so the size or scale of values does not affect how the tree makes its decision.

In []:

In []: